

Choosing count models

Class activity

https://sta214-s23.github.io/class_activities/ca_lecture_24.html

Goal: How can we use quantile residual plots to determine whether to fit a Poisson, quasi-Poisson, or negative binomial model?

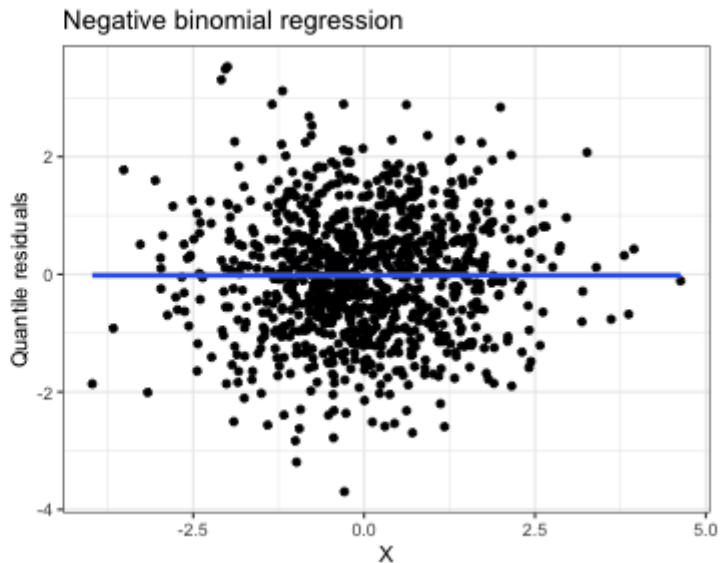
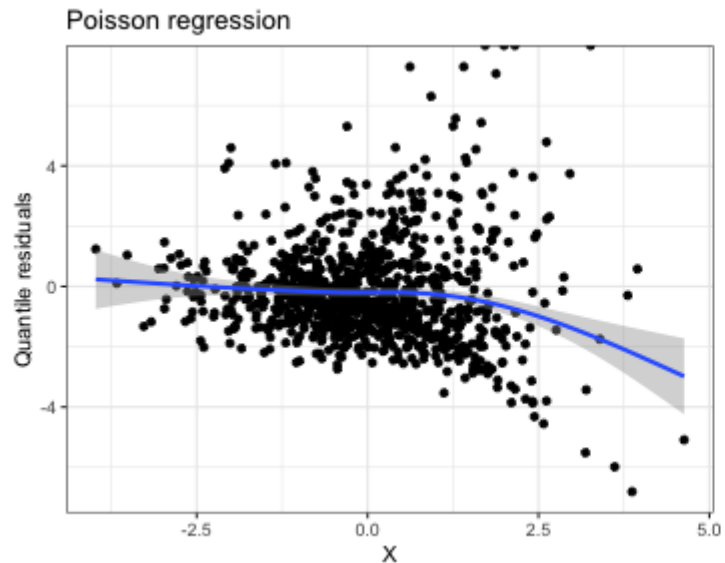
Instructions:

- + With a neighbor, work through the class activity
- + You and your neighbor will then discuss the activity with another pair of students from the class
- + Finally, we will discuss the activity as a group

Proposed guidelines

Class activity

Quantile residual plots for negative binomial data ($\mu_i = \exp(X_i), \theta = 0.5$):



Class activity

Data: negative binomial with $\mu_i = \exp(X_i)$, $\theta = 0.5$

How does CI coverage compare for Poisson, quasi-Poisson, and negative binomial models?

Class activity

Data: negative binomial with $\mu_i = \exp(X_i)$, $\theta = 0.5$

Poisson coverage:

```
n <- 1000
nsim <- 500
theta <- 0.5
contains_beta_poisson <- rep(0, nsim)
for(i in 1:nsim){
  x <- rnorm(n, sd = 1.2)
  y <- rnbinom(n, size=theta, mu=exp(x))

  m1 <- glm(y ~ x, family = poisson)

  upper <- summary(m1)$coefficients[2,1] +
    1.96*summary(m1)$coefficients[2,2]
  lower <- summary(m1)$coefficients[2,1] -
    1.96*summary(m1)$coefficients[2,2]

  contains_beta_poisson[i] <- upper > 1 && lower < 1
}
mean(contains_beta_poisson)
```

```
## [1] 0.278
```

Class activity

Data: negative binomial with $\mu_i = \exp(X_i)$, $\theta = 0.5$

quasi-Poisson coverage:

```
contains_beta_quasipoisson <- rep(0, nsim)
for(i in 1:nsim){
  x <- rnorm(n, sd = 1.2)
  y <- rnbinom(n, size=theta, mu=exp(x))

  m2 <- glm(y ~ x, family = quasipoisson)

  upper <- summary(m2)$coefficients[2,1] +
    qt(0.025, n-2, lower.tail=F)*summary(m2)$coefficients[2,2]
  lower <- summary(m2)$coefficients[2,1] -
    qt(0.025, n-2, lower.tail=F)*summary(m2)$coefficients[2,2]

  contains_beta_quasipoisson[i] <- upper > 1 && lower < 1
}
mean(contains_beta_quasipoisson)
```

```
## [1] 0.586
```

Class activity

Data: negative binomial with $\mu_i = \exp(X_i)$, $\theta = 0.5$

negative binomial coverage:

```
contains_beta_nb <- rep(0, nsim)
for(i in 1:nsim){
  x <- rnorm(n, sd = 1.2)
  y <- rnbinom(n, size=theta, mu=exp(x))

  m3 <- glm.nb(y ~ x)

  upper <- summary(m3)$coefficients[2,1] +
    1.96*summary(m3)$coefficients[2,2]
  lower <- summary(m3)$coefficients[2,1] -
    1.96*summary(m3)$coefficients[2,2]

  contains_beta_nb[i] <- upper > 1 && lower < 1
}
mean(contains_beta_nb)
```



```
## [1] 0.95
```


Class activity

What happens to quantile residual plots as I increase θ ? What about coverage?

Class activity

Quantile residual plots for negative binomial data ($\mu_i = \exp(X_i), \theta = 2$):



Class activity

Data: negative binomial with $\mu_i = \exp(X_i)$, $\theta = 2$

Poisson coverage:

```
n <- 1000
nsim <- 500
theta <- 2
contains_beta_poisson <- rep(0, nsim)
for(i in 1:nsim){
  x <- rnorm(n, sd = 1.2)
  y <- rnbinom(n, size=theta, mu=exp(x))

  m1 <- glm(y ~ x, family = poisson)

  upper <- summary(m1)$coefficients[2,1] +
    1.96*summary(m1)$coefficients[2,2]
  lower <- summary(m1)$coefficients[2,1] -
    1.96*summary(m1)$coefficients[2,2]

  contains_beta_poisson[i] <- upper > 1 && lower < 1
}
mean(contains_beta_poisson)
```

```
## [1] 0.508
```

Class activity

Data: negative binomial with $\mu_i = \exp(X_i)$, $\theta = 2$

quasi-Poisson coverage:

```
contains_beta_quasipoisson <- rep(0, nsim)
for(i in 1:nsim){
  x <- rnorm(n, sd = 1.2)
  y <- rnbinom(n, size=theta, mu=exp(x))

  m2 <- glm(y ~ x, family = quasipoisson)

  upper <- summary(m2)$coefficients[2,1] +
    qt(0.025, n-2, lower.tail=F)*summary(m2)$coefficients[2,2]
  lower <- summary(m2)$coefficients[2,1] -
    qt(0.025, n-2, lower.tail=F)*summary(m2)$coefficients[2,2]

  contains_beta_quasipoisson[i] <- upper > 1 && lower < 1
}
mean(contains_beta_quasipoisson)
```

```
## [1] 0.644
```

Class activity

Data: negative binomial with $\mu_i = \exp(X_i)$, $\theta = 2$

negative binomial coverage:

```
contains_beta_nb <- rep(0, nsim)
for(i in 1:nsim){
  x <- rnorm(n, sd = 1.2)
  y <- rnbinom(n, size=theta, mu=exp(x))

  m3 <- glm.nb(y ~ x)



  upper <- summary(m3)$coefficients[2,1] +
    1.96*summary(m3)$coefficients[2,2]
  lower <- summary(m3)$coefficients[2,1] -
    1.96*summary(m3)$coefficients[2,2]

  contains_beta_nb[i] <- upper > 1 && lower < 1
}
mean(contains_beta_nb)
```

```
## [1] 0.946
```

Class activity

Now let's simulate quasi-Poisson data with $\phi = 3$:



Class activity

Data: quasi-Poisson with $\phi = 3$

Poisson coverage:

```
n <- 1000
nsim <- 500
phi <- 3
contains_beta_poisson <- rep(0, nsim)
for(i in 1:nsim){
  x <- rnorm(n, sd = 1.2)
  y <- rqpois(n, mean = exp(x), dispersion = phi)

  m1 <- glm(y ~ x, family = poisson)

  upper <- summary(m1)$coefficients[2,1] +
    1.96*summary(m1)$coefficients[2,2]
  lower <- summary(m1)$coefficients[2,1] -
    1.96*summary(m1)$coefficients[2,2]

  contains_beta_poisson[i] <- upper > 1 && lower < 1
}
mean(contains_beta_poisson)
```

```
## [1] 0.738
```

Class activity

Data: quasi-Poisson with $\phi = 3$

quasi-Poisson coverage:

```
contains_beta_quasipoisson <- rep(0, nsim)
for(i in 1:nsim){
  x <- rnorm(n, sd = 1.2)
  y <- y <- rqpois(n, mean = exp(x), dispersion = phi)

  m2 <- glm(y ~ x, family = quasipoisson)

  upper <- summary(m2)$coefficients[2,1] +
    qt(0.025, n-2, lower.tail=F)*summary(m2)$coefficients[2,2]
  lower <- summary(m2)$coefficients[2,1] -
    qt(0.025, n-2, lower.tail=F)*summary(m2)$coefficients[2,2]

  contains_beta_quasipoisson[i] <- upper > 1 && lower < 1
}
mean(contains_beta_quasipoisson)
```

```
## [1] 0.92
```


Class activity

Data: quasi-Poisson with $\phi = 3$

negative binomial coverage:

```
contains_beta_nb <- rep(0, nsim)
for(i in 1:nsim){
  x <- rnorm(n, sd = 1.2)
  y <- y <- rqpois(n, mean = exp(x), dispersion = phi)

  m3 <- glm.nb(y ~ x)

  upper <- summary(m3)$coefficients[2,1] +
    1.96*summary(m3)$coefficients[2,2]
  lower <- summary(m3)$coefficients[2,1] -
    1.96*summary(m3)$coefficients[2,2]

  contains_beta_nb[i] <- upper > 1 && lower < 1
}
mean(contains_beta_nb)
```

```
## [1] 0.936
```

Choosing a count model with quantile residual plots

- + If the residuals have **constant variance** and mostly fall between **-2 and 2**: Poisson is reasonable
- + If the residuals have **constant variance** but many residuals are > 2 or < -2 : use either quasi-Poisson or negative binomial
- + If the residuals have **non-constant variance**: use negative binomial

quasi-Poisson vs. negative binomial

quasi-Poisson:

- + linear relationship between mean and variance
- + easy to interpret $\hat{\phi}$
- + same as Poisson regression when $\phi = 1$
- + simple adjustment to estimated standard errors
- + estimated coefficients same as in Poisson regression
- + t -tests and F -tests

negative binomial:

- + quadratic relationship between mean and variance
- + we get to use a likelihood, rather than a quasi-likelihood
- + Same as Poisson regression when θ is very large and p is very small
- + Wald tests and likelihood ratio tests